Hadoop in Windows Server using PolyBase

Requirements

Required Software: (Hadoop/ Hadoop bin files/Hadoop Configuration files/Java)

https://drive.google.com/open?id=1e6pYQxZrr1JQaAR5ywOWNxkEYjQKKpBL

SQL Server 2017 Developer Edition https://www.microsoft.com/en-us/sql-server/sql-server-downloads

SQL Server Edition **Standard, Web, Express with Advanced Services, Express** does not support head node they require head node

Enterprise Edition only can be used for head node and other can be Scale out with multiple compute nodes.

Before we start Some of the difference between Hadoop and SQL

HADOOP	SQL
Schema on Read Approach	**Schema on Write Approach**
When we write the data on Hadoop, has distributed file system so brings all the files without any rules. Then when we dictate the data we apply rules to the code that reads the data rather of preconfiguring the data ahead of time.	In Schema on Write while coping the data from database A to B the data is first checked before inserting to the database (checked the datatypes of the data and the data size) if not meet the requirement reject the data
Data Storage Process Hadoop	**Data Storage Process in SQL**
Data is stored in Compressed file of either text or others datatype. In the moment files or data stored are replicated in multiple nodes in Hadoop distributed filing system	Data is stored in a Logical form in Related Table and Columns

Installation of Java

Before installing Hadoop File Distributed System (HDFS) and **PolyBase** feature in SQL Server in 2016 or later requires Java. We need to first install Java in our system.

PolyBase enables your SQL Server 2016 and later instance to process Transact-SQL queries that read data from Hadoop. The same query can also access relational tables in your SQL Server. **PolyBase** enables the same query to also join the data from Hadoop and SQL Server. **PolyBase** is a mediator between SQL Server and Hadoop.

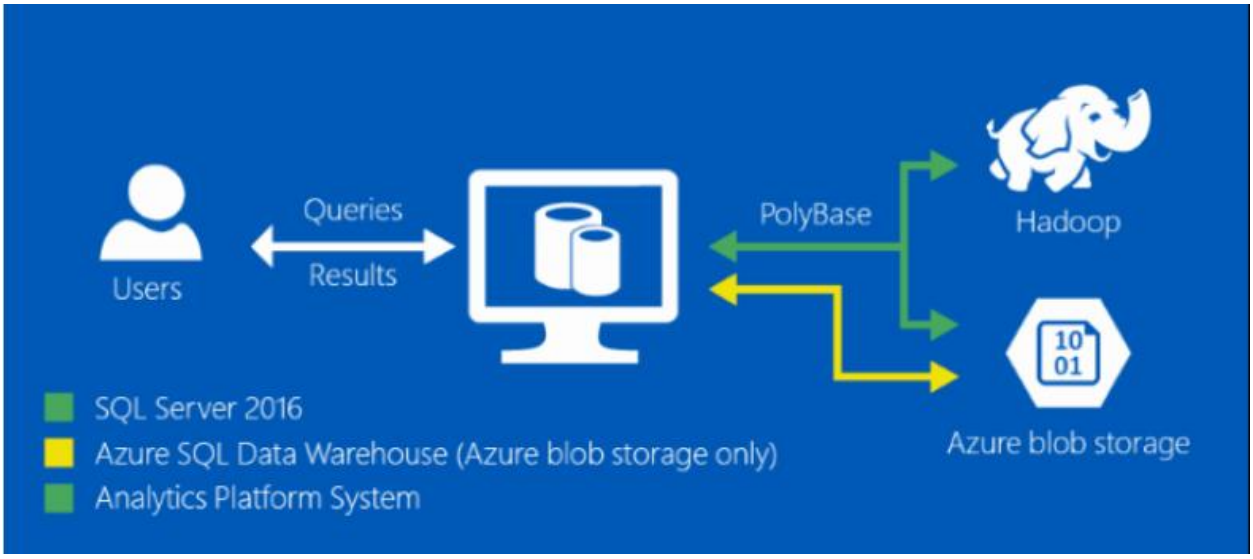

Set up

1. Check either Java 1.8.0 is already installed on your system or not, use **"Java -version"** to check.

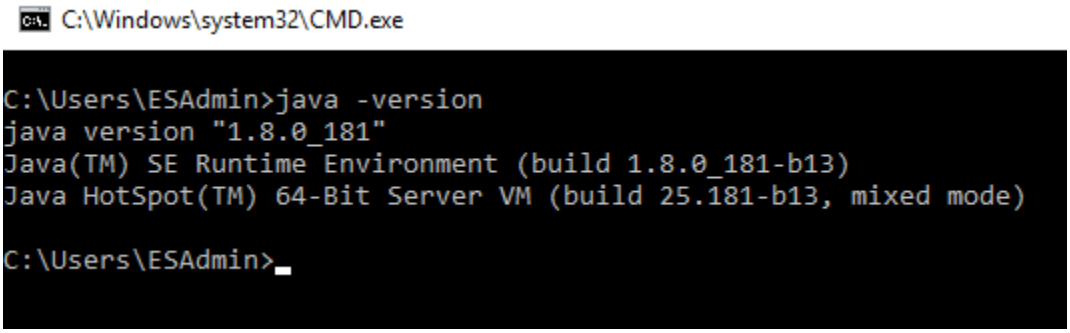

2. If Java is not installed on your system then first install java under **"C:\JAVA"**

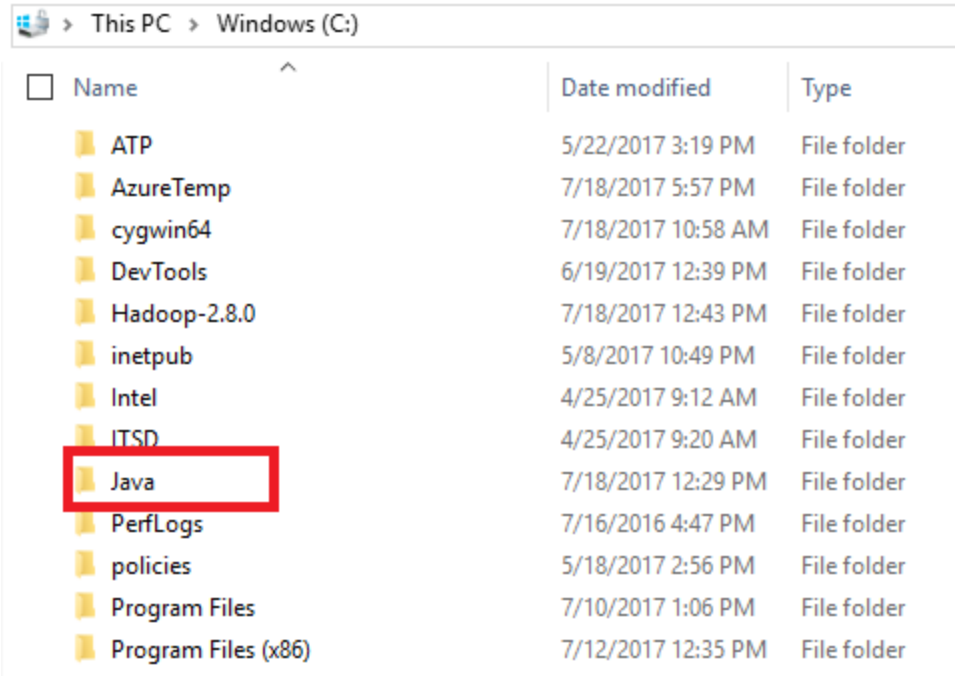

3. Extract file Hadoop 2.8.0.tar.gz or Hadoop-2.8.0.zip and place under **"C:\Hadoop-2.8.0"**.

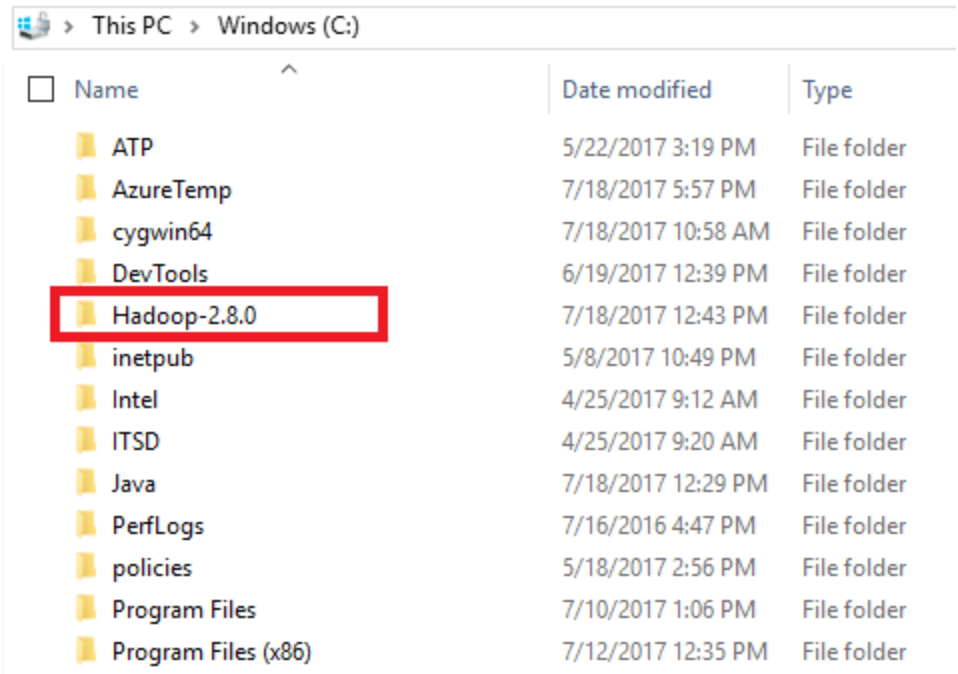

4. Set the path HADOOP_HOME Environment variable on windows 10(see Step 1,2,3 and 4 below).

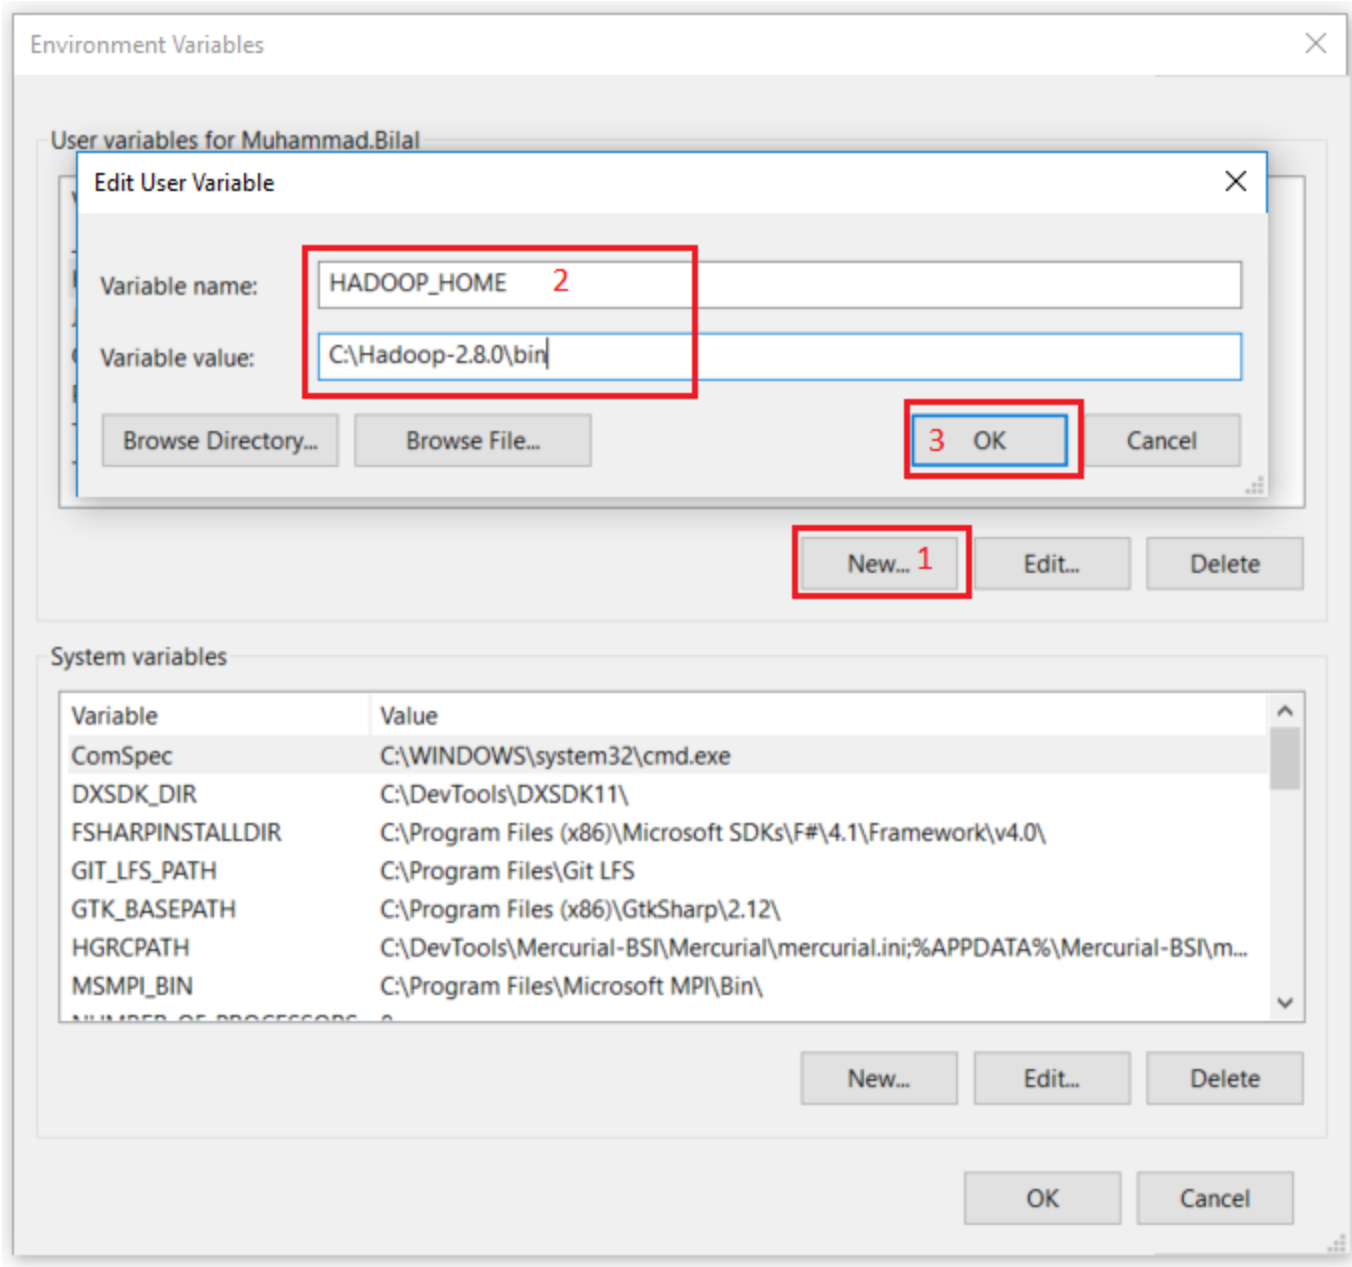

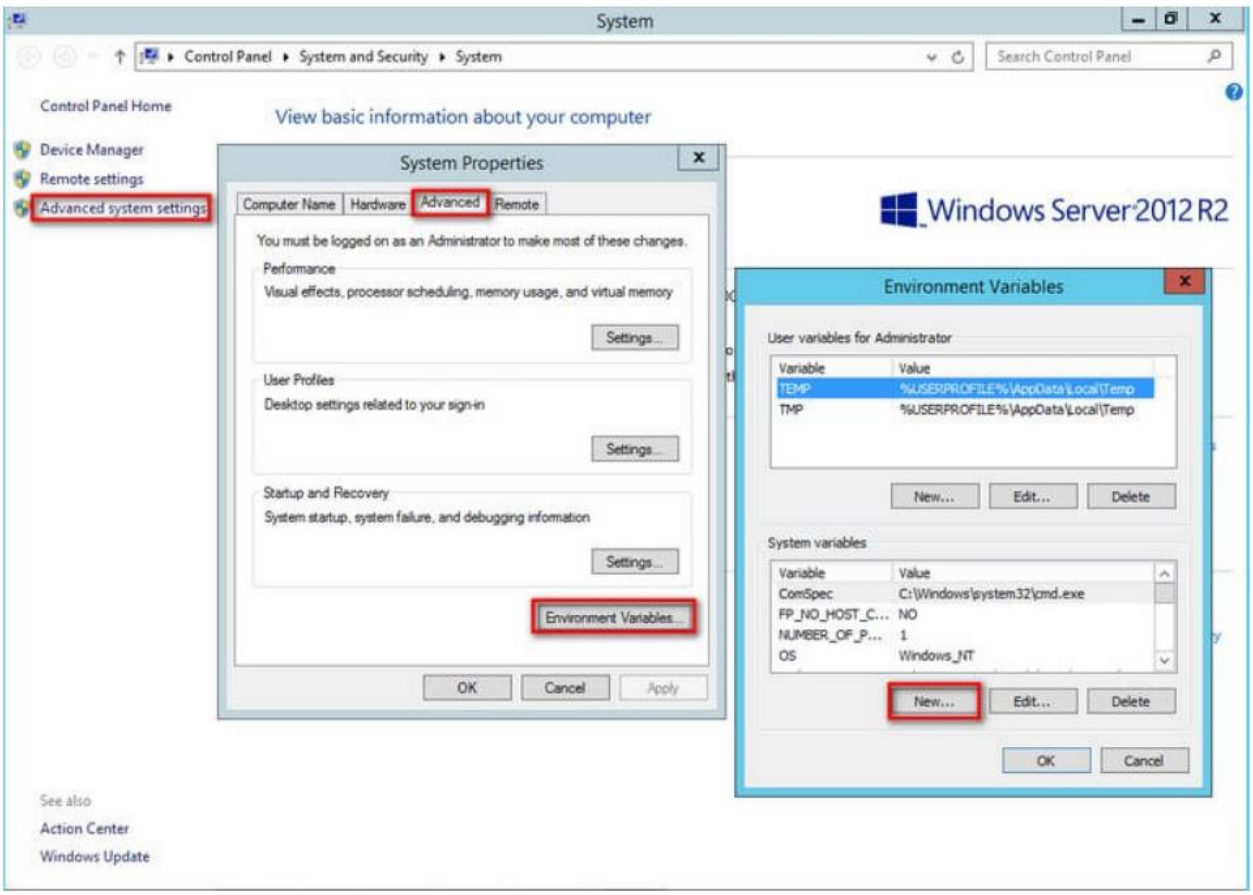
System
Control Panel ▸ System and Security ▸ System
Search Control Panel
Control Panel Home
View basic information about your computer
Device Manager
Remote settings
Advanced system settings
Windows Server 2012 R2
System Properties
Computer Name
Hardware
Advanced
Remote
You must be logged on as an Administrator to make most of these changes.
Performance
Visual effects, processor scheduling, memory usage, and virtual memory
Settings...
User Profiles
Desktop settings related to your sign-in
Settings...
Startup and Recovery
System startup, system failure, and debugging information
Settings...
Environment Variables...
OK
Cancel
Apply
Environment Variables
User variables for Administrator
Variable
Value
TEMP
%USERPROFILE%\AppData\Local\Temp
TMP
%USERPROFILE%\AppData\Local\Temp
New...
Edit...
Delete
System variables
Variable
Value
ComSpec
C:\Windows\system32\cmd.exe
FP_NO_HOST_C...
NO
NUMBER_OF_P...
1
OS
Windows_NT
New...
Edit...
Delete
OK
Cancel
See also
Action Center
Windows Update

5. Set the path JAVA_HOME Environment variable on windows 10(see Step 1,2,3 and 4 below).

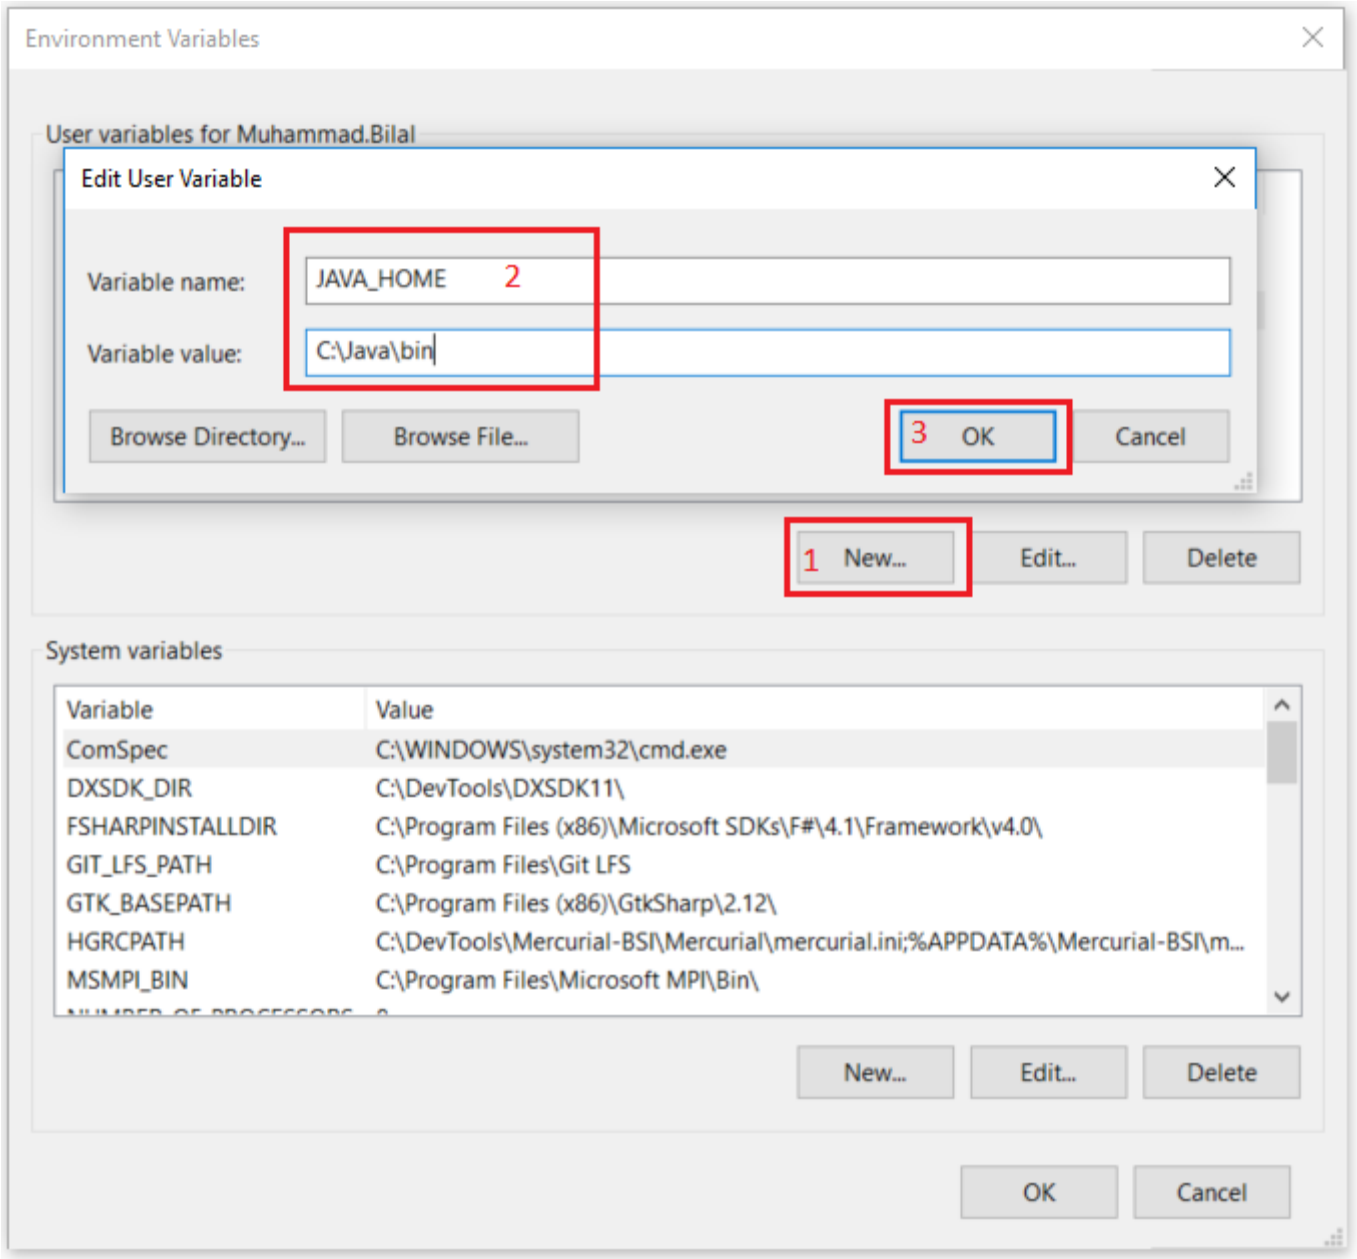

6. Next, we set the Hadoop bin directory path and JAVA bin directory path.

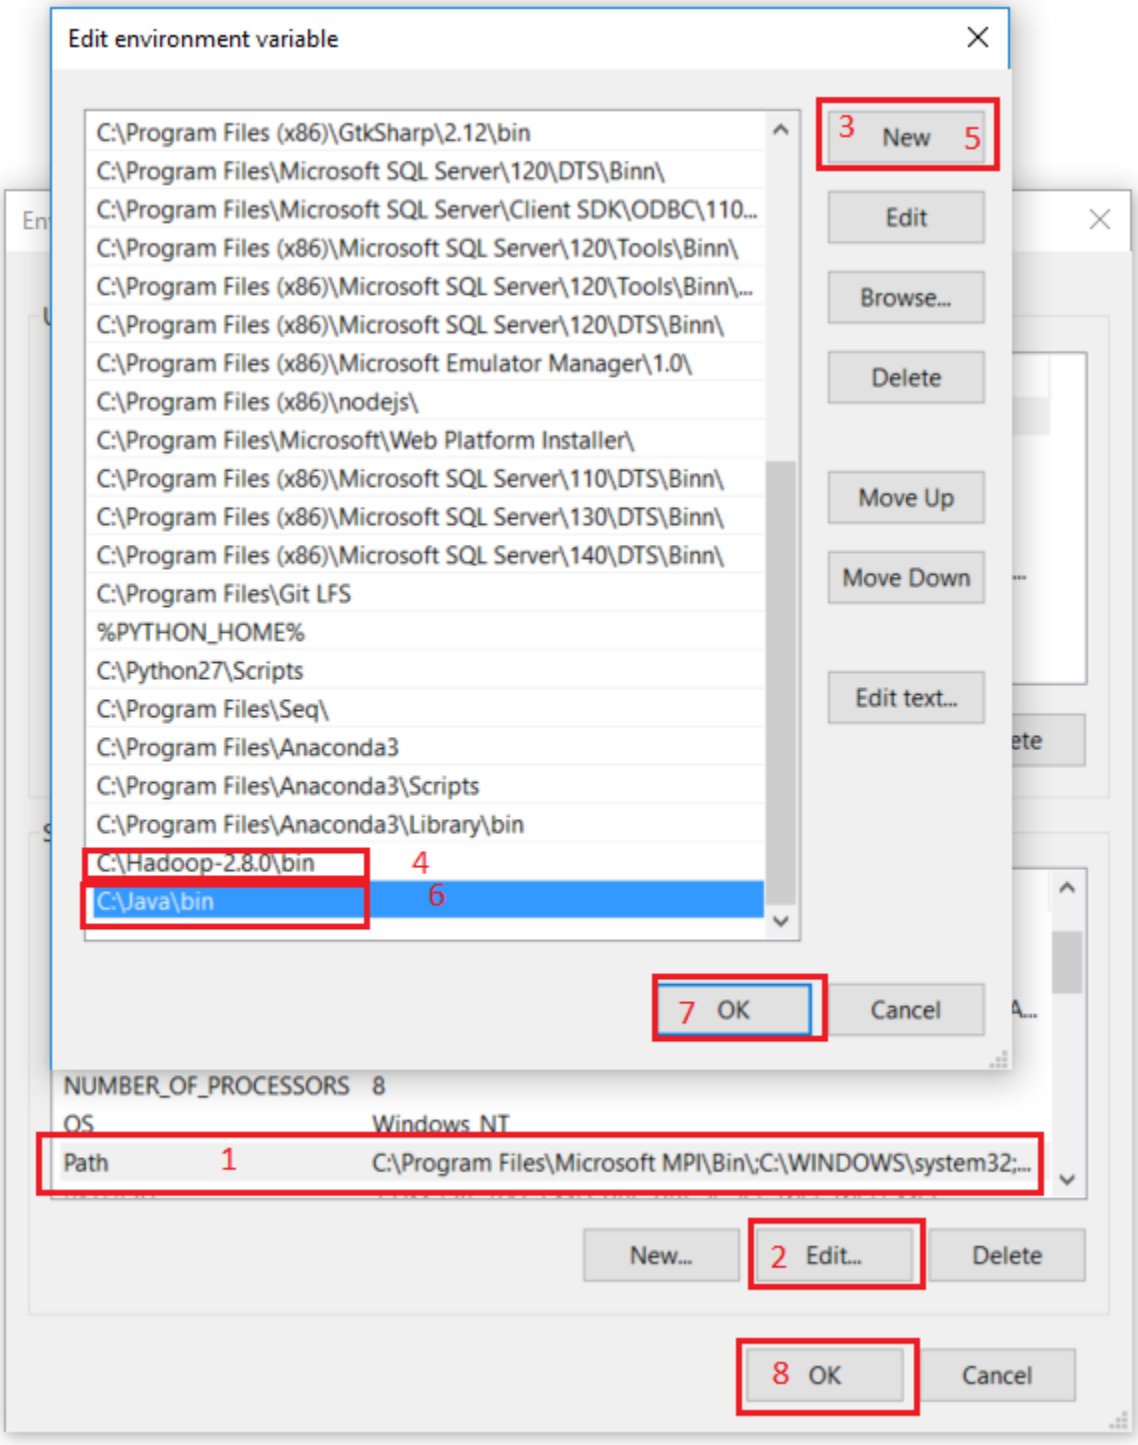

Install **Polybase** SQL Server 2016 or Later

Since **Polybase** is now part of SQL Server, we can use the SQL Server 2016 installation media to do the installation. And because it was designed to interact with Hadoop, we will need to install the Oracle Java SE Runtime Environment (JRE) 7.51 (x64) or higher prior to running the SQL Server 2016 installation media.

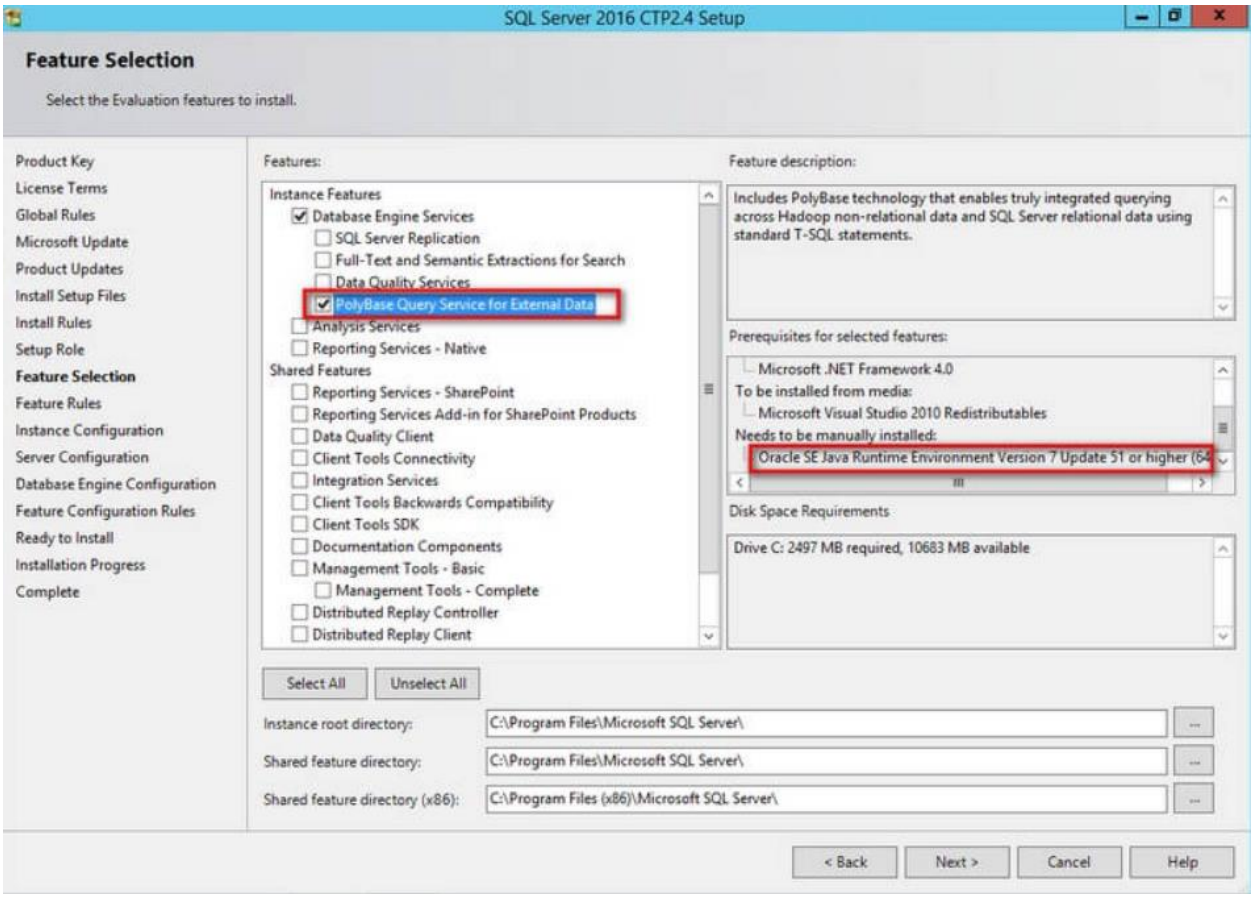

If You don't have JRE installed then the installation will fail and screen as below appears.

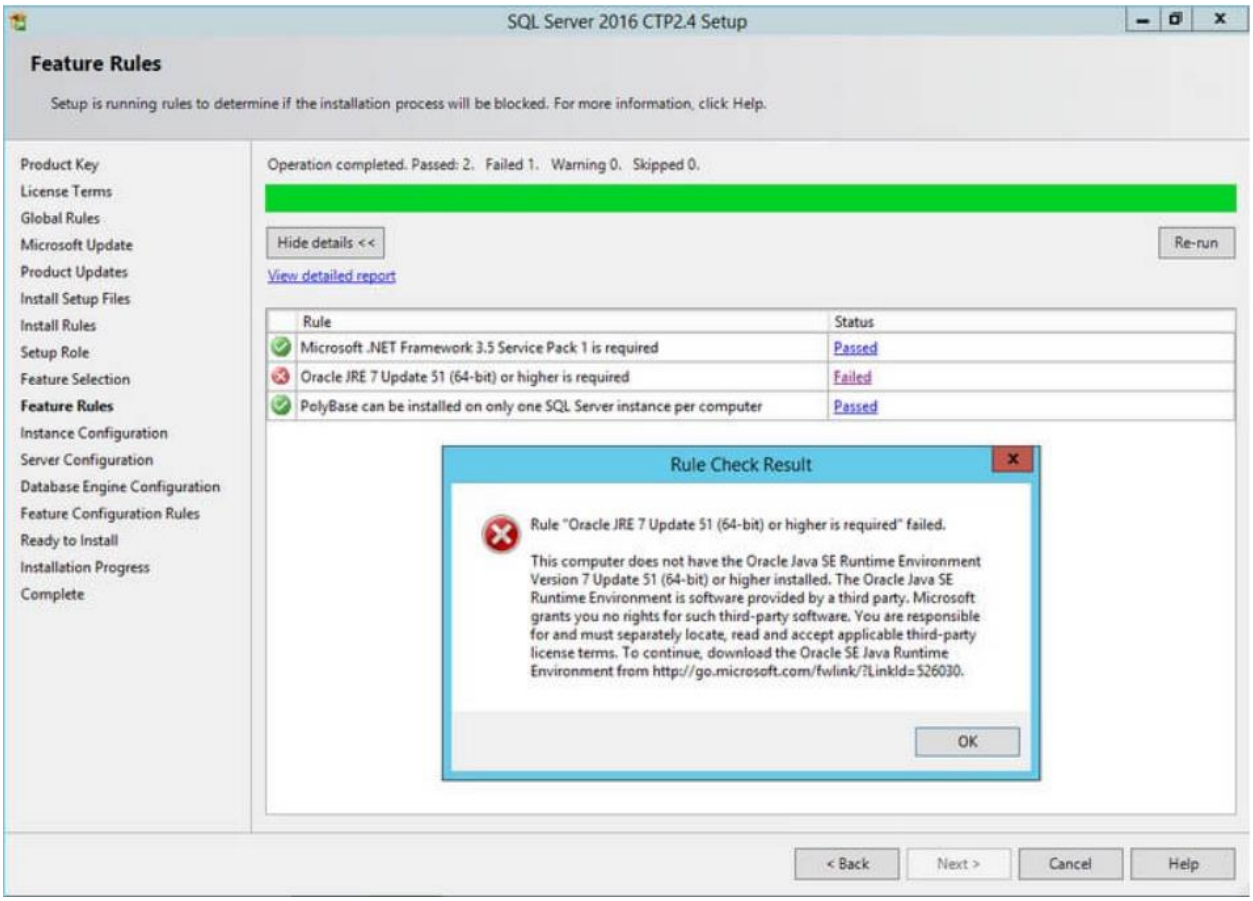
SQL Server 2016 CTP2.4 Setup
Feature Rules
Setup is running rules to determine if the installation process will be blocked. For more information, click Help.
Product Key
License Terms
Global Rules
Microsoft Update
Product Updates
Install Setup Files
Install Rules
Setup Role
Feature Selection
Feature Rules
Instance Configuration
Server Configuration
Database Engine Configuration
Feature Configuration Rules
Ready to Install
Installation Progress
Complete
Operation completed. Passed: 2. Failed 1. Warning 0. Skipped 0.
Hide details <<
Re-run
View detailed report
Rule
Status
Microsoft .NET Framework 3.5 Service Pack 1 is required
Passed
Oracle JRE 7 Update 51 (64-bit) or higher is required
Failed
PolyBase can be installed on only one SQL Server instance per computer
Passed
Rule Check Result
Rule "Oracle JRE 7 Update 51 (64-bit) or higher is required" failed.
This computer does not have the Oracle Java SE Runtime Environment Version 7 Update 51 (64-bit) or higher installed. The Oracle Java SE Runtime Environment is software provided by a third party. Microsoft grants you no rights for such third-party software. You are responsible for and must separately locate, read and accept applicable third-party license terms. To continue, download the Oracle SE Java Runtime Environment from http://go.microsoft.com/fwlink/?LinkId=526030.
OK
< Back
Next >
Cancel
Help

Install JRE and set the **Environment Variable** after installation which is mention in **Set Up** steps above.

And Continue the installation of Polybase and completed successfully.

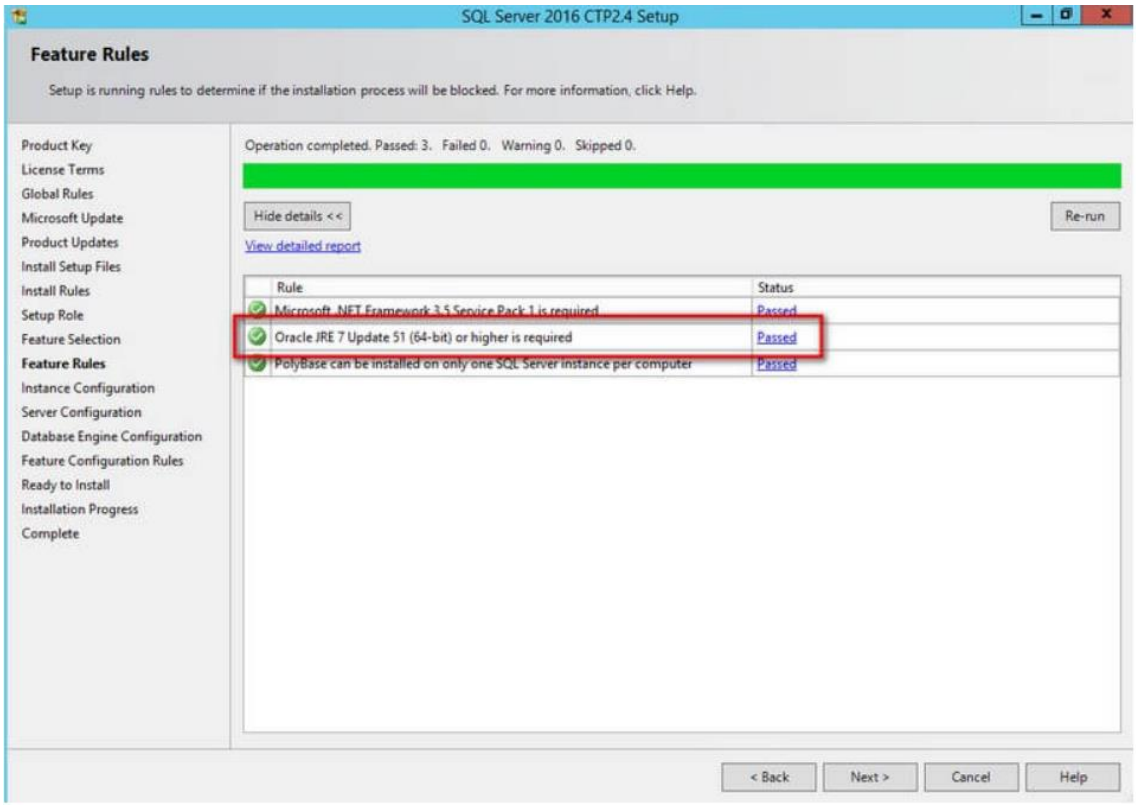

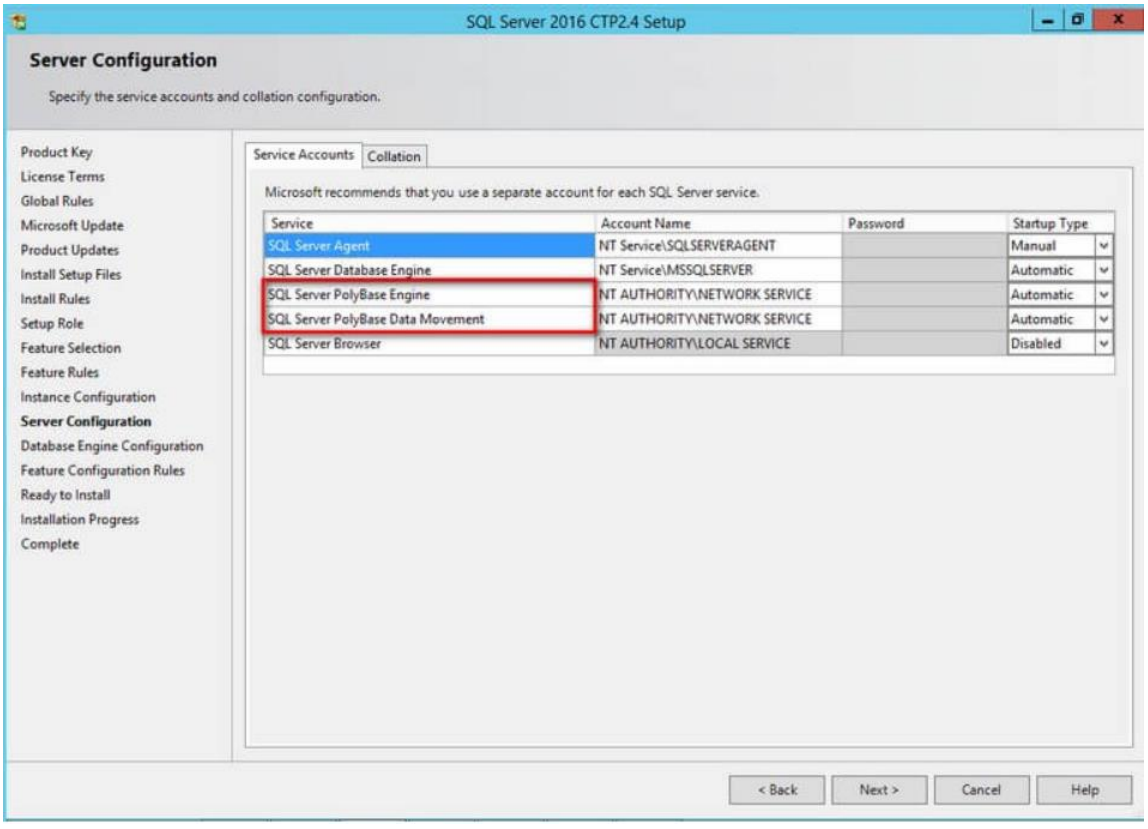

Configuration

1. Edit file **C:/Hadoop-2.8.0/etc/hadoop/core-site.xml**, paste below xml paragraph and save this file.

```
<configuration>
   <property>
        <name>fs.defaultFS</name>
        <value>hdfs://localhost:9000</value>
   </property>
</configuration>
```

2. Rename "mapred-site.xml.template" to "mapred-site.xml" and edit this file **C:/Hadoop-2.8.0/etc/hadoop/mapred-site.xml**, paste below xml paragraph and save this file.

```
<configuration>
   <property>
        <name>mapreduce.framework.name</name>
        <value>yarn</value>
   </property>
</configuration>
```

3. Create folder **"data"** under **"C:\Hadoop-2.8.0"**

- Create folder **"datanode"** under **"C:\Hadoop-2.8.0\data"**
- Create folder **"namenode"** under **"C:\Hadoop-2.8.0\data"**

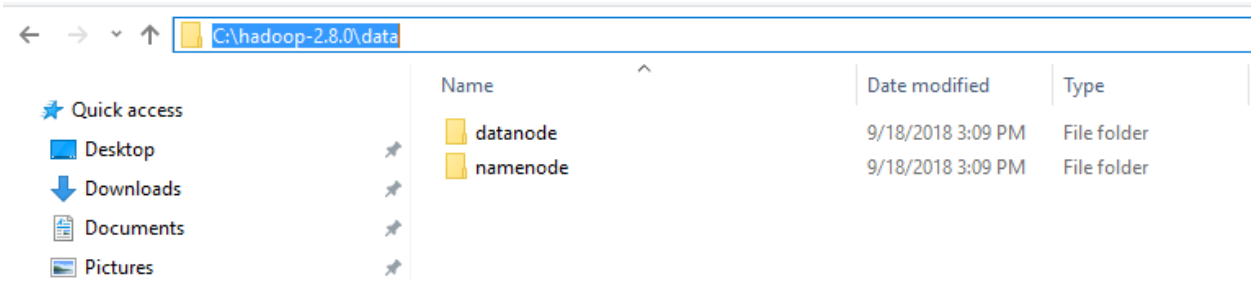

4. Edit file **C:\Hadoop-2.8.0/etc/hadoop/hdfs-site.xml**, paste below xml paragraph and save this file.

```
<configuration>
   <property>
        <name>dfs.replication</name>
        <value>1</value>
   </property>
   <property>
        <name>dfs.namenode.name.dir</name>
        <value>C:\hadoop-2.8.0\data\namenode</value>
   </property>
   <property>
        <name>dfs.datanode.data.dir</name>
        <value>C:\hadoop-2.8.0\data\datanode</value>
   </property>
</configuration>
```

5. Edit file **C:/Hadoop-2.8.0/etc/hadoop/yarn-site.xml**, paste below xml paragraph and save this file.

```
<configuration>
   <property>
         <name>yarn.nodemanager.aux-services</name>
         <value>mapreduce_shuffle</value>
   </property>
   <property>
         <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
         <value>org.apache.hadoop.mapred.ShuffleHandler</value>
   </property>
</configuration>
```

6. Edit file **C:/Hadoop-2.8.0/etc/hadoop/hadoop-env.cmd** by closing the command line**"JAVA_HOME=%JAVA_HOME%"** instead of set **"JAVA_HOME=C:\Java"** (On C:\java this is path to file jdk.18.0)

```
@rem The java implementation to use.  Required.
@rem set JAVA_HOME=%JAVA_HOME%
set JAVA_HOME=C:\Java\jdk1.8.0_181
@rem The jsvc implementation to use. Jsvc is required to run secure datanodes.
@rem set JSVC_HOME=%JSVC_HOME%
```

Hadoop Configuration

1. Download file Hadoop Configuration.zip (Link: https://drive.google.com/open?id=1e6pYQxZrr1JQaAR5ywOWNxkEYjQKKpBL)
2. Delete file bin on C:\Hadoop-2.8.0\bin, replaced by file bin on file just download (from Hadoop Configuration.zip).
3. Open cmd and typing command **"hdfs namenode –format"**. You will see

Testing

1. Open cmd and change directory to "C:\Hadoop-2.8.0\sbin" and type **"start-all.cmd"** to start
apache.

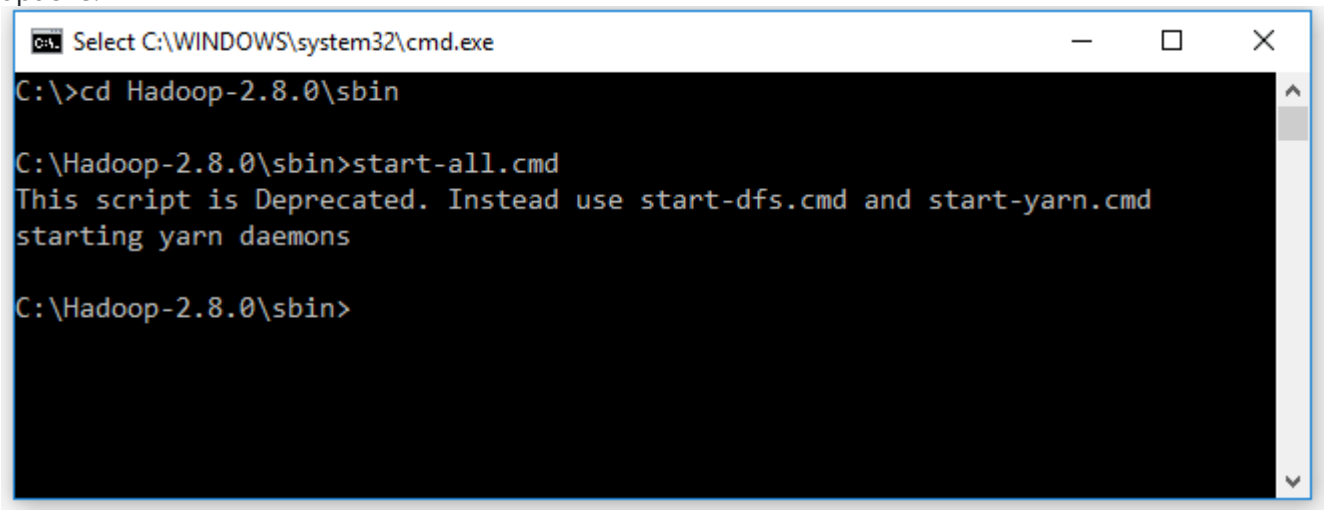

2. Make sure these apps are running

- Hadoop Namenode
- Hadoop datanode
- YARN Resource Manager
- YARN Node Manager

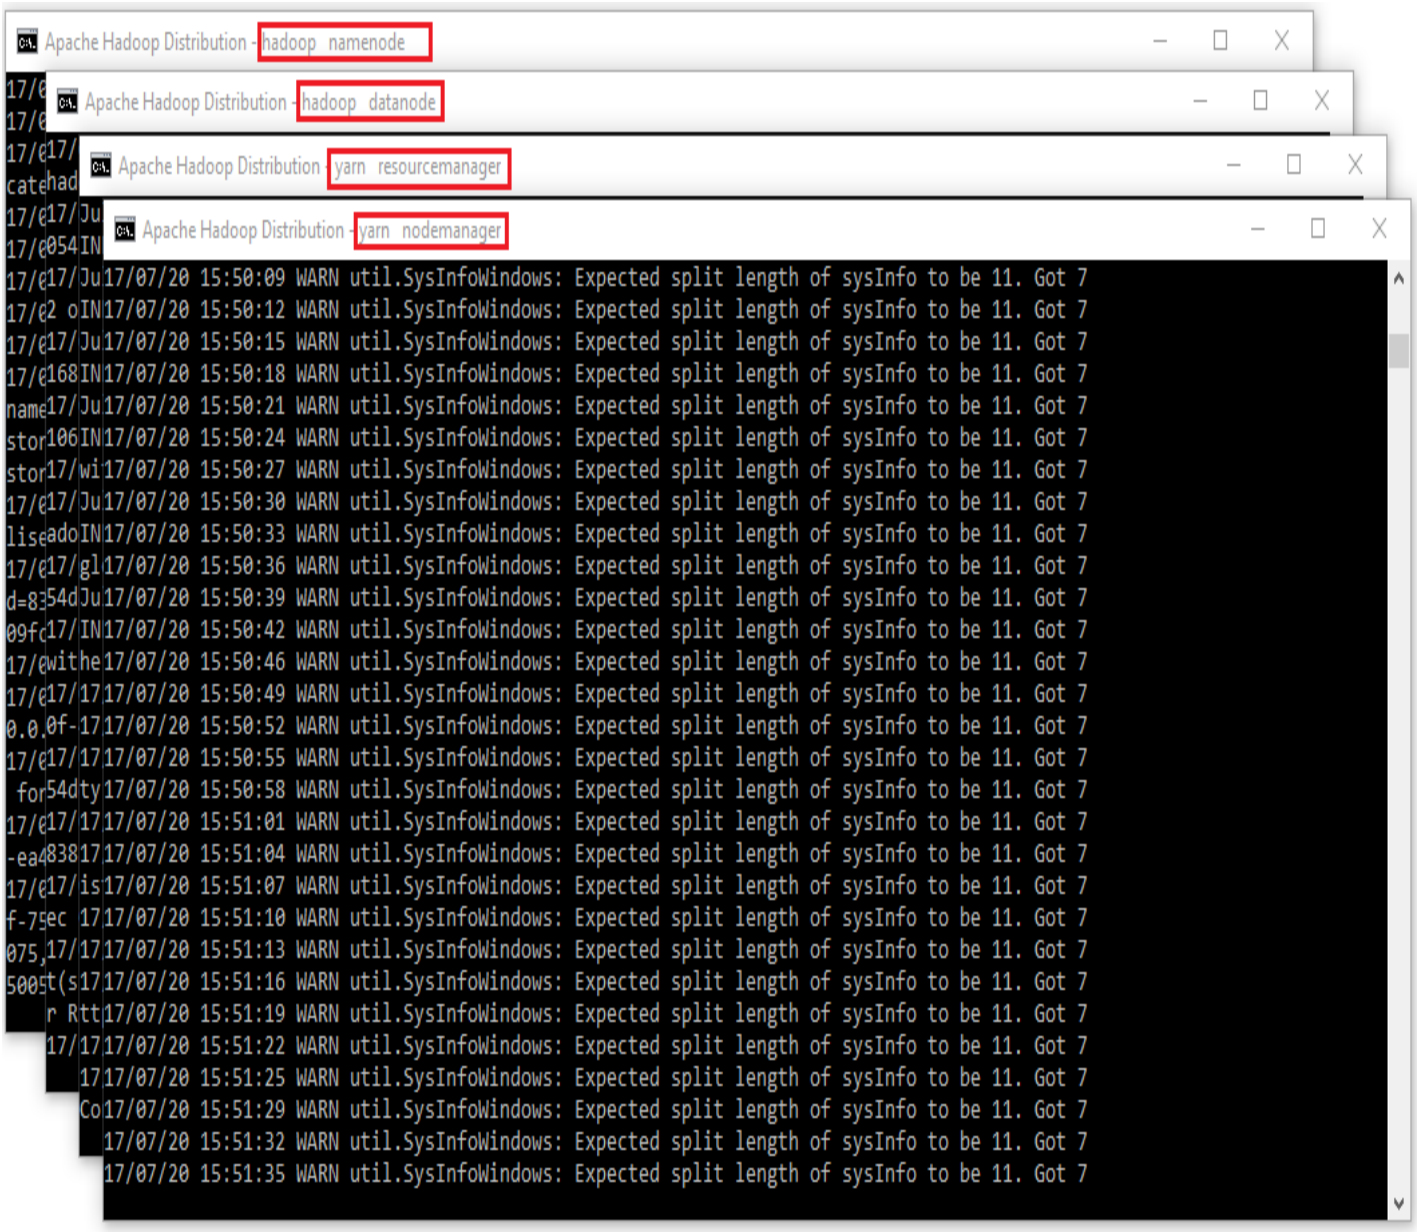
Apache Hadoop Distribution - hadoop namenode
Apache Hadoop Distribution - hadoop datanode
Apache Hadoop Distribution - yarn resourcemanager
Apache Hadoop Distribution - yarn nodemanager
17/07/20 15:50:09 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:12 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:15 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:18 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:21 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:24 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:27 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:30 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:33 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:36 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:39 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:42 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:46 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:49 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:52 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:55 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:50:58 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:01 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:04 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:07 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:10 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:13 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:16 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:19 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:22 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:25 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:29 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:32 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
17/07/20 15:51:35 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7

Open: http://localhost:8088

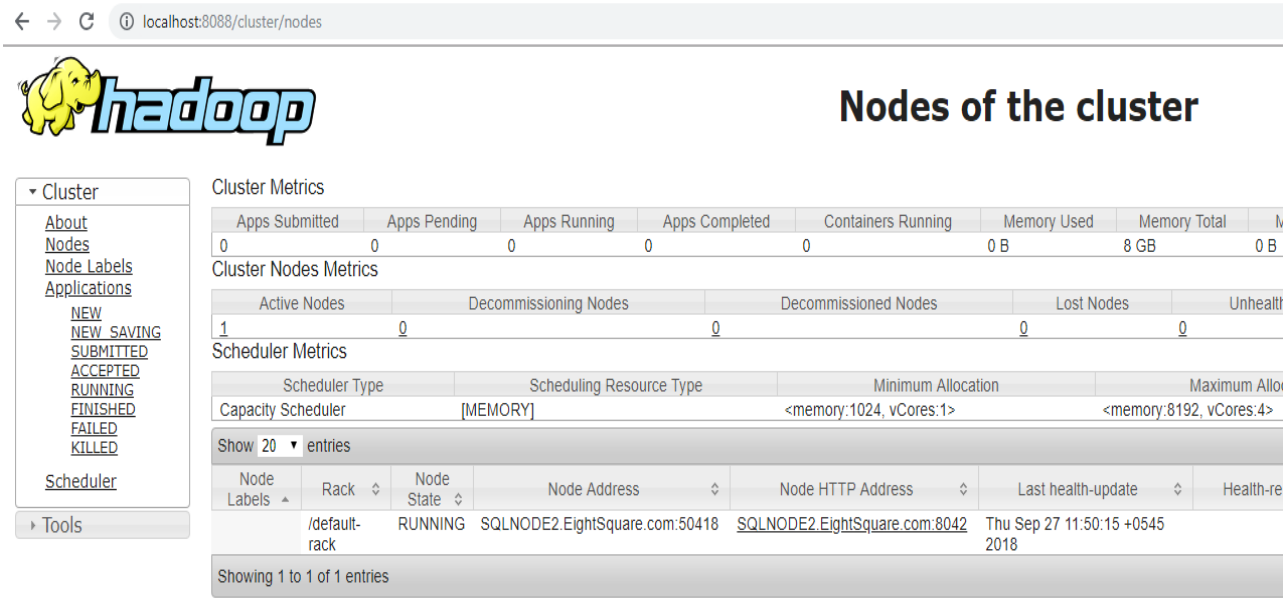

Open: http://localhost:50070

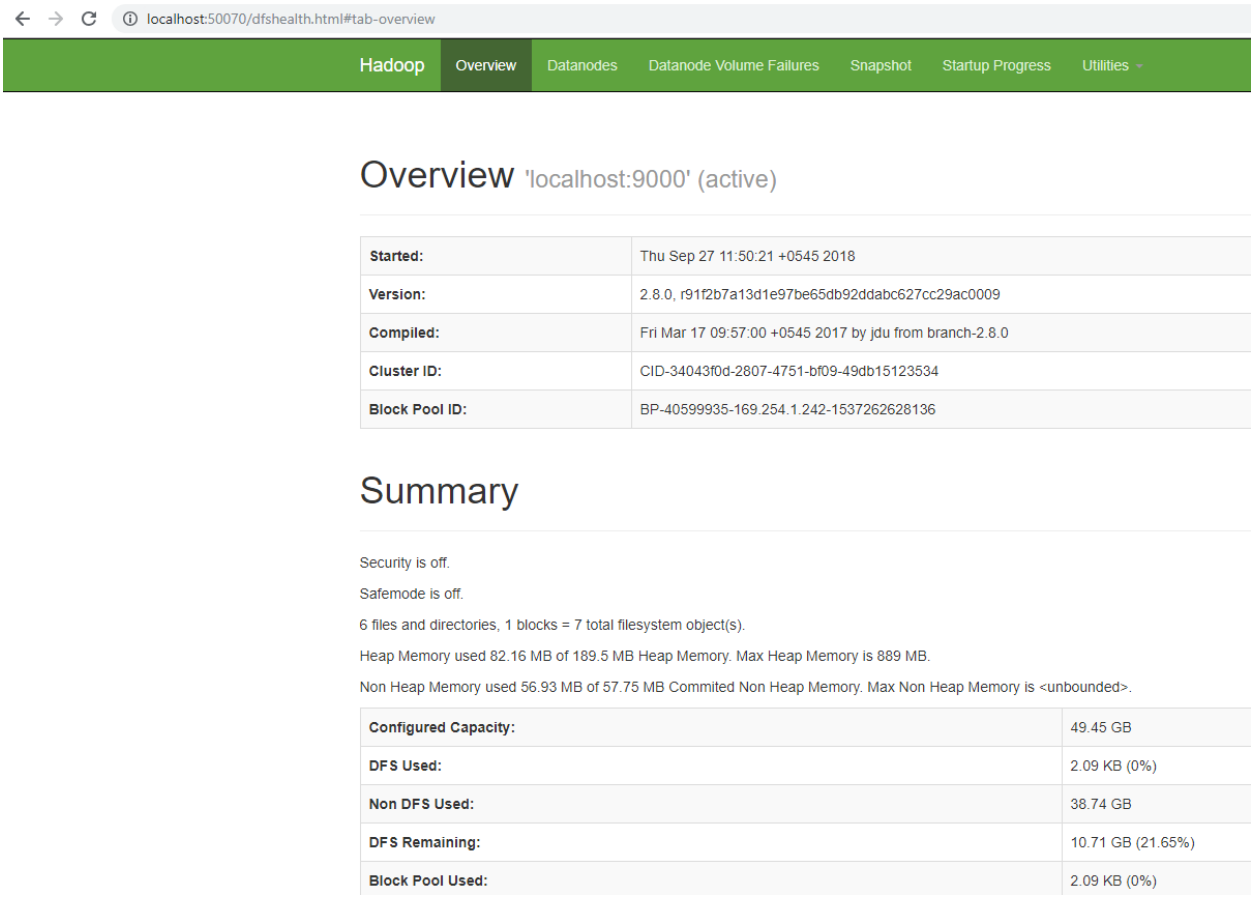

Create Directory in Hadoop

Import file to MMESDataFiles directory

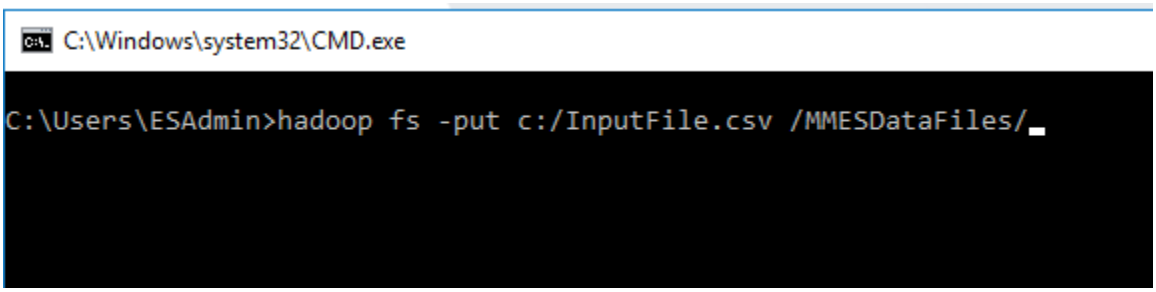

Simple procedure to configure database for **PolyBase**, Hadoop Connectivity,
And Create External Data Source, Create External File Format and Create External Table also create Statistics on external table for query optimization.

```
CREATE DATABASE PolybaseDB
use PolybaseDB

SELECT SERVERPROPERTY ('IsPolybaseInstalled') AS IsPolybaseInstalled;

--Prestep: Configuring Hadoop flavor
exec sp_configure 'hadoop connectivity',7
Reconfigure

exec sp_configure 'allow polybase export',1
Reconfigure

--(a) Creating external data source --Hadoop  HDP Cluster

--DROP EXTERNAL DATA SOURCE [HadoopCluster]
CREATE EXTERNAL DATA SOURCE HadoopCluster
WITH (TYPE = Hadoop,
      LOCATION = N'hdfs://localhost:9000')

select * from sys.external_data_sources;

--(b) Creating external file formates -delimited text

--DROP EXTERNAL FILE FORMAT [TextFile]
CREATE EXTERNAL FILE FORMAT TextFile
WITH (FORMAT_TYPE = DelimitedText,
      FORMAT_OPTIONS (FIELD_TERMINATOR = N',',
      USE_TYPE_DEFAULT = True));


select * from sys.external_file_formats;
```

```
-- (c) Creating external tables refering to data in external Hadoop Cluster

--DROP EXTERNAL TABLE [dbo].[SensorDataHDP]
CREATE EXTERNAL TABLE [dbo].[SensorDataHDP]
(
    Id int NOT NULL,
    BeneAccountCreditId int NOT NULL,
    BeneficiaryId int NULL,
    SendAmount varchar(50) NOT NULL,
    TypeofTraxId int NOT NULL
)
WITH (LOCATION = '/MMESDataFiles/',
      DATA_SOURCE = HadoopCluster,
      FILE_FORMAT = TextFile,
      REJECT_TYPE = Value,
      REJECT_VALUE = 0
);

SELECT * FROM SensorDataHDP
INSERT INTO SensorDataHDP
SELECT
      [Id]
     ,[BeneAccountCreditId]
     ,[BeneficiaryId]
     ,[SendAmount]
     ,[TypeofTraxId]
FROM [PolybaseDB].[dbo].[EremitData]
where Id < 100

SELECT * FROM SensorDataHDP

CREATE STATISTICS STAT_SensorDataHDPSendAmount on SensorDataHDP(SendAmount)

SELECT
      hdp.*
FROM SensorDataHDP hdp
INNER JOIN [EremitData] ed on hdp.Id = ed.Id
WHERE hdp.SendAmount between 100 and 200
```

```
SELECT
      hdp.*
FROM SensorDataHDP hdp
INNER JOIN [EremitData] ed on hdp.Id = ed.Id
WHERE hdp.SendAmount between 100 and 200
```

150 %

Results | Messages

	Id	BeneAccountCreditId	BeneficiaryId	SendAmount	TypeofTraxId
1	3	4035	4035	125	1
2	44	827	827	100	1
3	50	827	827	100	1
4	82	536	536	190	1

Structure of External Data Source, External File Format and External Table in SSMS.

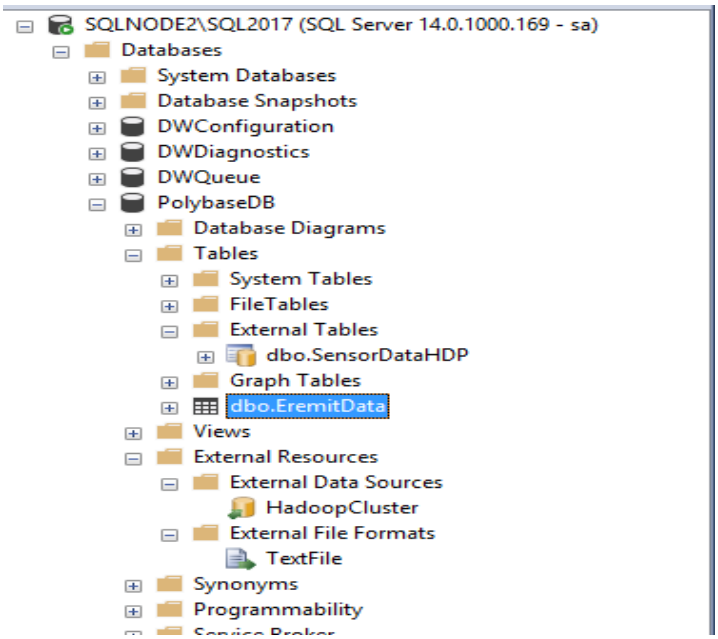

File Resides in HDFS